

Predictive_Analytics_Project_1

Subhayan Biswas

January 21, 2026

0.1 Problem Set 1: An Introduction

In this script, we will be working with the ‘Boston’ housing data which originates from MASS Library of R.

```
[ ]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
[ ]: df = pd.read_csv('Boston.csv')  
df
```

```
[ ]:      Unnamed: 0      crim      zn  indus  chas      nox      rm      age      dis      rad  \  
0            1  0.00632  18.0    2.31      0  0.538  6.575  65.2  4.0900      1  
1            2  0.02731  0.0    7.07      0  0.469  6.421  78.9  4.9671      2  
2            3  0.02729  0.0    7.07      0  0.469  7.185  61.1  4.9671      2  
3            4  0.03237  0.0    2.18      0  0.458  6.998  45.8  6.0622      3  
4            5  0.06905  0.0    2.18      0  0.458  7.147  54.2  6.0622      3  
..          ...     ...     ...     ...     ...     ...     ...     ...     ...     ...  
501           502  0.06263  0.0   11.93      0  0.573  6.593  69.1  2.4786      1  
502           503  0.04527  0.0   11.93      0  0.573  6.120  76.7  2.2875      1  
503           504  0.06076  0.0   11.93      0  0.573  6.976  91.0  2.1675      1  
504           505  0.10959  0.0   11.93      0  0.573  6.794  89.3  2.3889      1  
505           506  0.04741  0.0   11.93      0  0.573  6.030  80.8  2.5050      1  
  
      tax  ptratio  black  lstat  medv  
0    296     15.3  396.90    4.98  24.0  
1    242     17.8  396.90    9.14  21.6  
2    242     17.8  392.83    4.03  34.7  
3    222     18.7  394.63    2.94  33.4  
4    222     18.7  396.90    5.33  36.2  
..          ...     ...     ...     ...  
501          273    21.0  391.99    9.67  22.4  
502          273    21.0  396.90    9.08  20.6  
503          273    21.0  396.90    5.64  23.9  
504          273    21.0  393.45    6.48  22.0  
505          273    21.0  396.90    7.88  11.9
```

[506 rows x 15 columns]

1. Report the ‘class’ of the data set. How many rows and columns are in this data set? What do the rows and columns represent?

[]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Unnamed: 0    506 non-null    int64  
 1   crim        506 non-null    float64 
 2   zn          506 non-null    float64 
 3   indus       506 non-null    float64 
 4   chas        506 non-null    int64  
 5   nox         506 non-null    float64 
 6   rm          506 non-null    float64 
 7   age         506 non-null    float64 
 8   dis          506 non-null    float64 
 9   rad          506 non-null    int64  
 10  tax          506 non-null    int64  
 11  ptratio     506 non-null    float64 
 12  black       506 non-null    float64 
 13  lstat       506 non-null    float64 
 14  medv        506 non-null    float64 
dtypes: float64(11), int64(4)
memory usage: 59.4 KB
```

From the output above, we can see that there are 15 columns and 506 entries or rows. Here, each row is a set of observations corresponding to the i th suburb or town. Whereas each of the columns refer to a distinct feature (variable). The columns are given as:

- **crim** - per capita crime rate by town
- **zn** - proportion of residential land zoned for lots over 25,000 sq.ft.
- **indus** - proportion of non-retail business acres per town.
- **chas** - Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).
- **nox** - nitrogen oxides concentration (parts per 10 million).
- **rm** - average number of rooms per dwelling.
- **age** - proportion of owner-occupied units built prior to 1940.
- **dis** - weighted mean of distances to five Boston employment centres.
- **rad** - index of accessibility to radial highways
- **tax** - full-value property-tax rate per \$10,000

- `ptratio` - pupil-teacher ratio by town.
 - `black` - $1000(Bk - 0.63)^2$ where `Bk` is the proportion of blacks by town.
 - `lstat` - lower status of the population (percent).
 - `medv` - median value of owner-occupied homes in \$1000s.
-

2. Create a smaller data set with the variables median value of owner-occupied homes, per capita crime rate, nitrogen oxides concentration, proportion of blacks and percentage of lower status of the population. Choosing median value of owner occupied homes as the response and the rest as the predictors, make scatter plots of the response versus each predictor. Present the scatter plots in different panels of the same graph. Comment on your findings.

```
[ ]: df_mini = pd.DataFrame(df,
                           columns = ['medv', 'crim', 'nox',
                                       'black', 'lstat'])
```

```
[ ]:      medv      crim      nox     black    lstat
0    24.0  0.00632  0.538  396.90   4.98
1    21.6  0.02731  0.469  396.90   9.14
2    34.7  0.02729  0.469  392.83   4.03
3    33.4  0.03237  0.458  394.63   2.94
4    36.2  0.06905  0.458  396.90   5.33
..
501   22.4  0.06263  0.573  391.99   9.67
502   20.6  0.04527  0.573  396.90   9.08
503   23.9  0.06076  0.573  396.90   5.64
504   22.0  0.10959  0.573  393.45   6.48
505   11.9  0.04741  0.573  396.90   7.88
```

[506 rows x 5 columns]

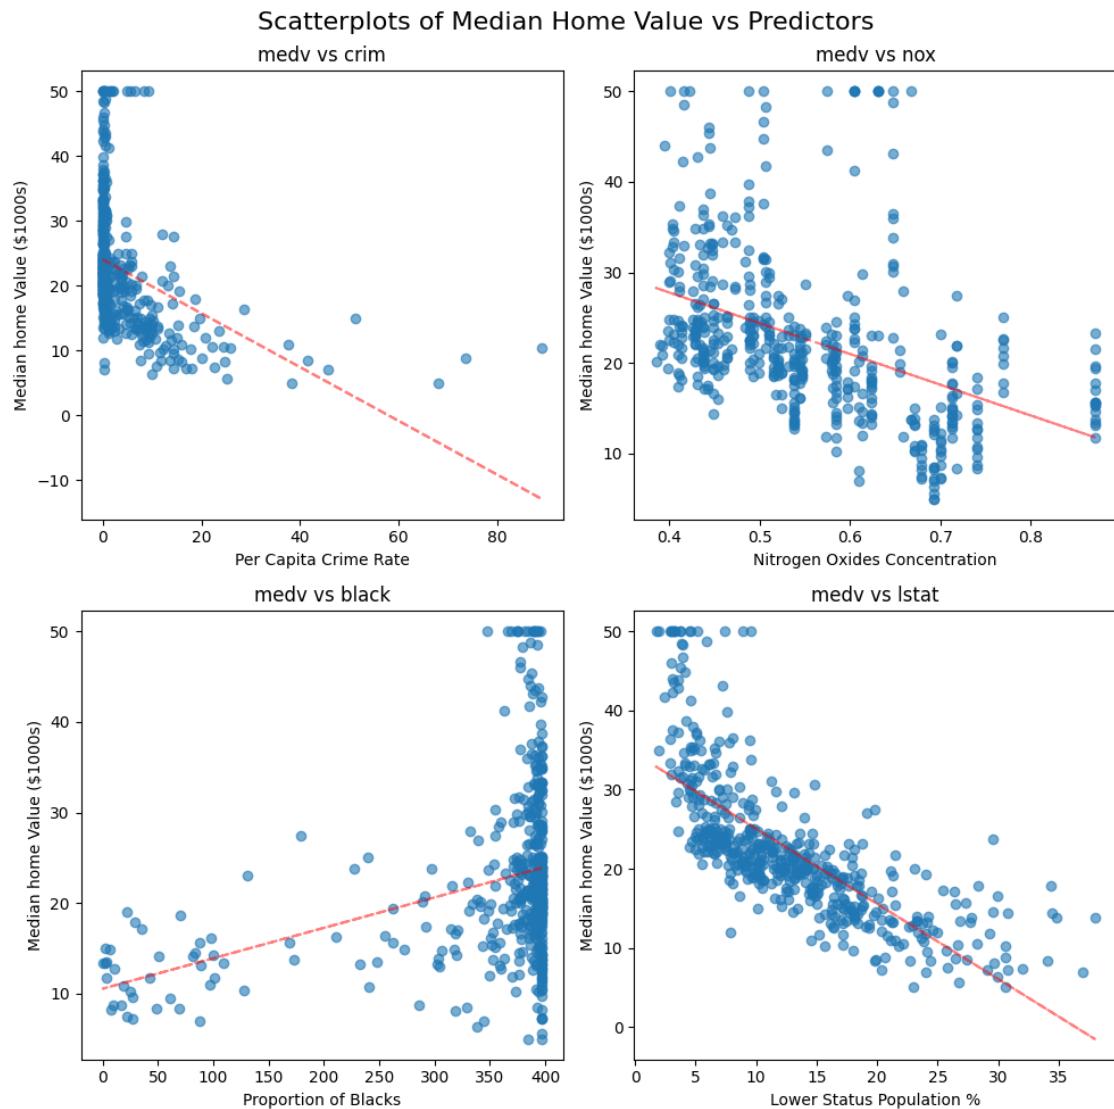
```
[ ]: fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(10, 10))
fig.suptitle('Scatterplots of Median Home Value vs Predictors',
             fontsize = 16)
predictors = ['crim', 'nox', 'black', 'lstat']
titles = ['Per Capita Crime Rate', 'Nitrogen Oxides Concentration',
          'Proportion of Blacks', 'Lower Status Population %']
for i, (predictor, title) in enumerate(zip(predictors, titles)):
    row, col = i//2, i % 2
    axes[row, col].scatter(df_mini[predictor],
                           df_mini['medv'], alpha = 0.6)
    axes[row, col].set_xlabel(title)
    axes[row, col].set_ylabel('Median home Value ($1000s)')
    axes[row, col].set_title(f'medv vs {predictor}')
```

```

z = np.polyfit(df_mini[predictor], df_mini['medv'], 1)
p = np.poly1d(z)
axes[row, col].plot(df_mini[predictor],
                     p(df_mini[predictor]), 'r--',
                     alpha = 0.5)

plt.tight_layout()
plt.show()

```



From the above scatter-plot matrix, we can see that: i) `medv` is negatively correlated with `crim` i.e. as crime rate increases, the median home value decreases and vice versa.

- ii) `medv` is moderately negatively correlated with `nox` i.e as the concentration of Nitrogen Oxides (a measure of air pollution) increases, home values generally decrease.

- iii) `medv` is somewhat positively correlated with `black`.
 - iv) `medv` is strongly negatively correlated with `lstat` suggesting that it might be one of the best predictors for the `medv`.
-

3. Which suburb of Boston has lowest median value of owner-occupied homes? What are the values of the other predictors mentioned in (2), for that suburb. How do these values compare to the overall ranges for those predictors? Comment on your findings. Hint: Mention which percentile these values belong to.

```
[ ]: df_mini[df_mini['medv'].min() == df_mini['medv']]
```

```
[ ]:      medv      crim      nox     black    lstat
398    5.0  38.3518  0.693   396.90  30.59
405    5.0  67.9208  0.693   384.97  22.98
```

```
[ ]: df_mini.describe()
```

```
[ ]:      medv      crim      nox     black    lstat
count  506.000000  506.000000  506.000000  506.000000  506.000000
mean   22.532806   3.613524   0.554695   356.674032  12.653063
std    9.197104   8.601545   0.115878   91.294864   7.141062
min    5.000000   0.006320   0.385000   0.320000   1.730000
25%   17.025000   0.082045   0.449000   375.377500  6.950000
50%   21.200000   0.256510   0.538000   391.440000  11.360000
75%   25.000000   3.677083   0.624000   396.225000  16.955000
max   50.000000   88.976200  0.871000   396.900000  37.970000
```

```
[ ]: from scipy.stats import percentileofscore
variables = ['crim', 'nox', 'black', 'lstat']
suburb_values = df_mini[df_mini['medv'].min() == df_mini['medv']]
percentiles = [percentileofscore(df[var], suburb_values[var],
                                  kind = 'rank')/100
               for var in variables]
percentiles
```

```
[ ]: [array([0.98814229, 0.99604743]),
      array([0.84486166, 0.84486166]),
      array([0.88142292, 0.34980237]),
      array([0.97826087, 0.89920949])]
```

In the light of our findings, we can see that there is not one but two suburbs that have the minimum value which is 5.00(in \$1000s). The values of the predictors corresponding to the minimum median value have also been shown in the output. Furthermore, we also derived the percentiles for each of the predictor values corresponding to the minimum median value and they are as follows:

- i) **Suburb-1**

- `crim`: 98.81%

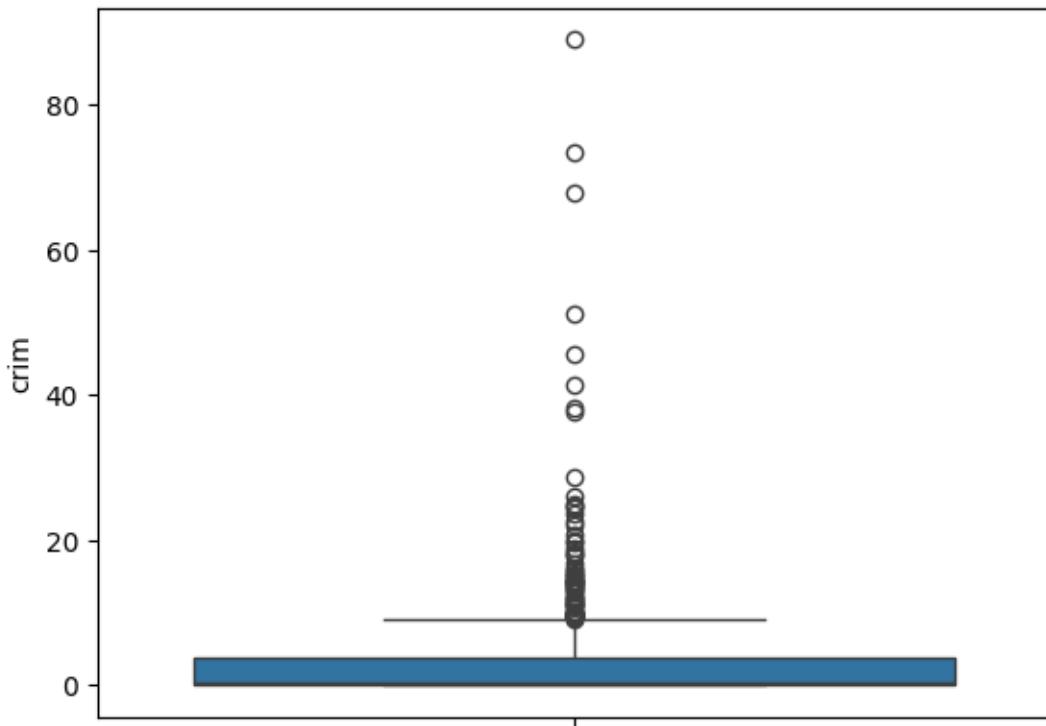
- nox: 84.49%
- black: 88.14%
- lstat: 97.93%

ii) Suburb-2

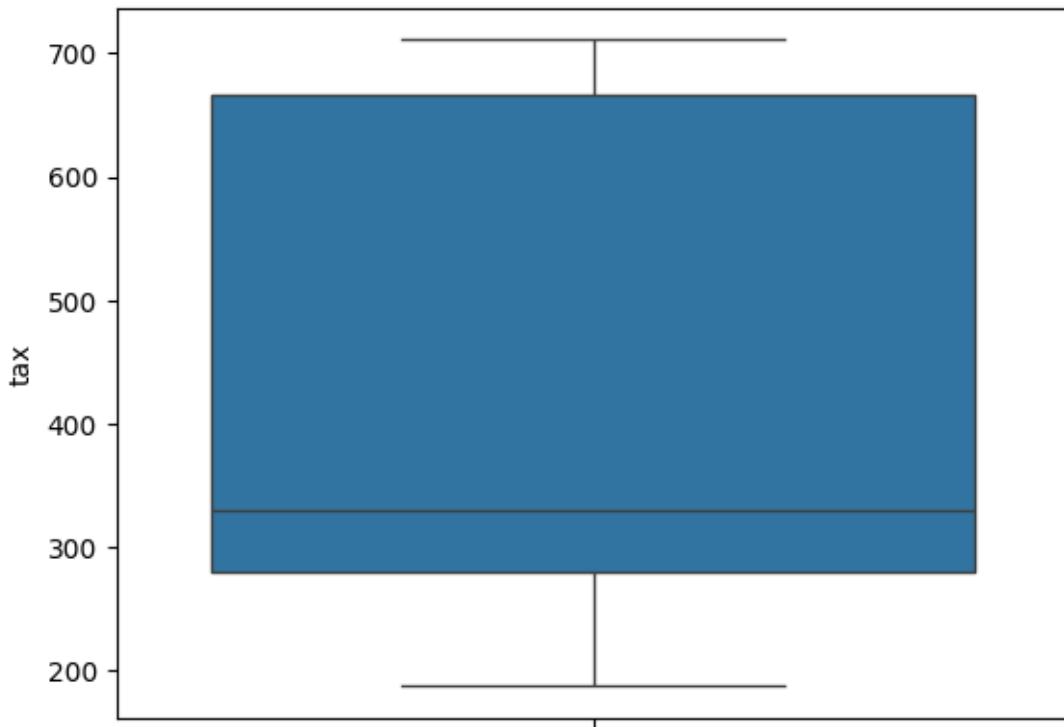
- crim: 99.60%
- nox: 84.49%
- black: 34.98%
- lstat: 89.92%

-
4. Does any suburb of Boston stand out for having notably high crime rates, tax rates, or pupil-teacher ratios? Hint: Use a boxplot to detect any outliers. If so, identify the suburbs that show the outlier values.

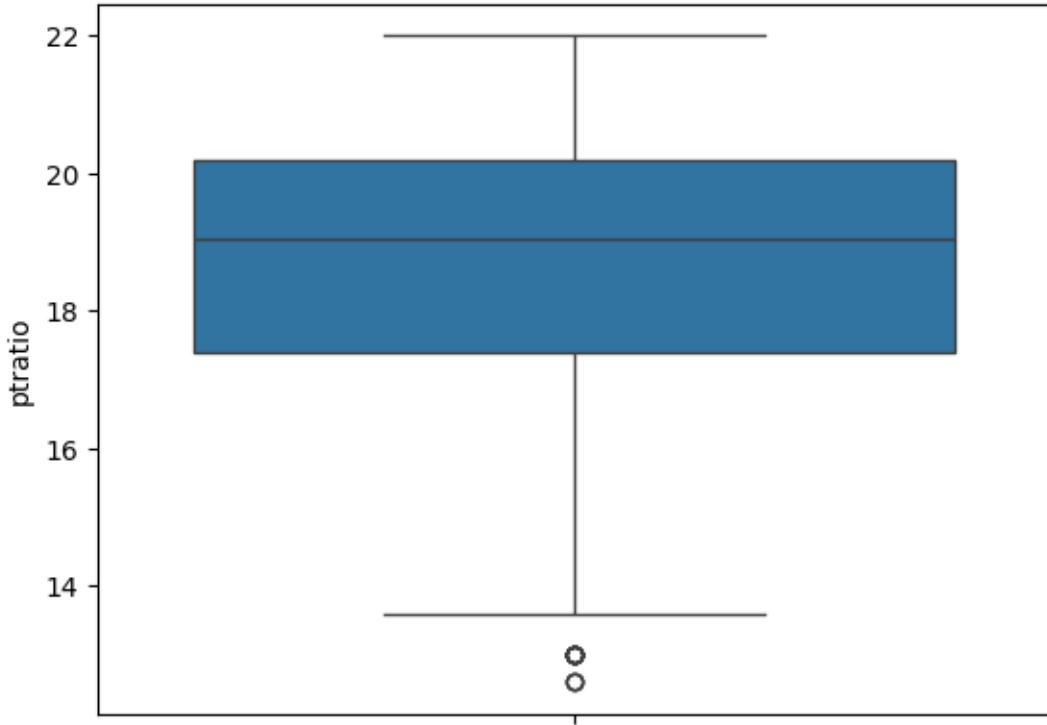
```
[ ]: import seaborn as sns  
sns.boxplot(data = df['crim'])  
plt.show()
```



```
[ ]: sns.boxplot(data = df['tax'])  
plt.show()
```



```
[ ]: sns.boxplot(data = df['ptratio'])
plt.show()
```



As we can see from the boxplots, some suburbs exhibit extremely high crime rates and some suburbs also show unusually high tax rates and pupil–teacher ratios. Hence, these outliers may strongly influence predictive models and should be handled carefully.
